# Information theory

- info should be high when an unlikely event has occured
- info should be zero when a guaranteed event ($P(X)=1$) has occured
- the information from 2 independent events should be additive.

$I(X=x) = -\log[P(X=x)]$ (captures these intuitions)

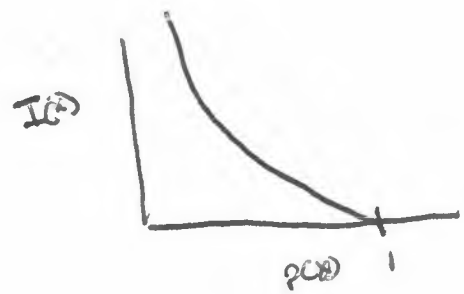


## entropy

- average amount of info in a distribution

eg:

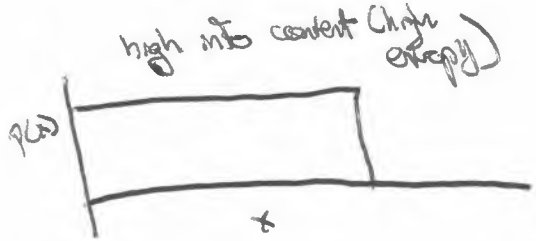


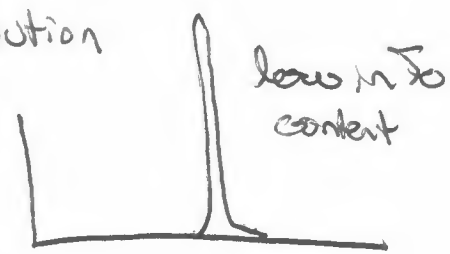


= avg. info contained w/in a dist: $H[X] = E_{x\sim p}[I[X]] = -E_{x\sim p}[P(X)]$

## KL divergence

$D_{KL}(P\|Q) = E_{x\sim p}\left[\frac{\log P(x)}{\log Q(x)}\right]$ ≈ a measure of difference between P and Q

(but, $D_{KL}(P\|Q) \neq D_{KL}(Q\|P)$)

- One can show that the MLE minimizes the KL divergence between the empirical distribution and the model distribution. This is another, nice, way of interpreting an MLE.